


JetNet-30 GPU-compute--quality Pareto comparison



Lower-left is better. GPU model FLOPs only; CPU preprocessing, metrics, I/O, and wall-clock time are excluded.

Conditional generation only; attribute/conditioning generators are excluded.

Training is the g/q/t portfolio cost: joint training is counted once and separate fits are summed. The training row shows one point per checkpoint.

Lifecycle adds inference for 10^6 generated jets. One GPU MAC = 2 FLOPs; operation precision is not throughput-weighted. Training uses published/released maximum budgets where stopping epochs are unavailable. JetFUEL guidance is selected separately by class.

The 128-step $w=0.25$ arm is pending; the w64/d4 ablation contributes only its nominated 64-step unguided gluon result.