


JetNet-30 GPU-compute-quality Pareto comparison (OmniLearn and EPiC-GAN omitted)



Lower-left is better. GPU model FLOPs only; CPU preprocessing, metrics, I/O, and wall-clock time are excluded. Conditional generation only; attribute/conditioning generators are excluded. Training is the portfolio cost for all applicable g/q/t models: joint training is counted once and separate fits are summed. The training row shows one point per checkpoint, using its best measured sampler for each class. Lifecycle adds inference for 10^7 total generated jets. One GPU MAC = 2 FLOPs; FP32 and FP64 operations are not throughput-weighted. Training totals use published/released maximum budgets where stopping epochs are unavailable. JetFUEL guidance points use the lowest measured FPND separately by class; the 128-step $w = 0.25$ arm is pending.