


JetNet-30 GPU-compute--quality Pareto comparison (OmniLearn and EPiC-GAN omitted)



Lower-left is better. GPU model FLOPs only; CPU preprocessing, metrics, I/O, and wall-clock time are excluded. Conditional generation only; attribute/conditioning generators are excluded. Training is the portfolio cost for all applicable g/q/t models: joint training is counted once and separate fits are summed. The training row shows one point per checkpoint, using its best measured sampler for each class. Lifecycle adds inference for  $10^8$  total generated jets. One GPU MAC = 2 FLOPs; FP32 and FP64 operations are not throughput-weighted. Training totals use published/released maximum budgets where stopping epochs are unavailable. JetFUEL guidance points use the lowest measured FPND separately by class; the 128-step  $w = 0.25$  arm is pending.